

Applied Capstone project - Car Accident Severity

September 19, 2020

1 Introduction/Business Problem

In an effort to reduce the frequency of car collisions in a community, an algorithm must be developed to predict the severity of an accident given the current weather, road and visibility conditions. When conditions are bad, this model will alert drivers to remind them to be more careful.

2 Data Description

Our predictor or target variable will be 'SEVERITYCODE' because it is used measure the severity of an accident from 0 to 5 within the dataset. Attributes used to weigh the severity of an accident are 'WEATHER', 'ROADCOND' and 'LIGHTCOND'.

Severity codes are as follows:

1. Little to no Probability (Clear Conditions)
2. Very Low Probability - Chance or Property Damage
3. Low Probability - Chance of Injury
4. Mild Probability - Chance of Serious Injury
5. High Probability - Chance of Fatality

Now we will import the necessary libraries & extract the dataset

```
[58]: import os
import numpy as np
import pandas as pd
from sklearn.utils import resample
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import jaccard_similarity_score
from sklearn.metrics import f1_score
from sklearn.metrics import log_loss
```

```
[4]: os.chdir(r'C:\Users\avira\OneDrive\Desktop\Extra Notes')
```

```
[5]: data = pd.read_csv('Data-Collisions.csv')
```

```
C:\Users\avira\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3058: DtypeWarning: Columns (33) have
```

mixed types. Specify dtype option on import or set low_memory=False.
interactivity=interactivity, compiler=compiler, result=result)

[6]: data

```
[6]:
```

	SEVERITYCODE	X	Y	OBJECTID	INCKEY	COLDETKEY	\
0	2	-122.323148	47.703140	1	1307	1307	
1	1	-122.347294	47.647172	2	52200	52200	
2	1	-122.334540	47.607871	3	26700	26700	
3	1	-122.334803	47.604803	4	1144	1144	
4	2	-122.306426	47.545739	5	17700	17700	
...	
194668	2	-122.290826	47.565408	219543	309534	310814	
194669	1	-122.344526	47.690924	219544	309085	310365	
194670	2	-122.306689	47.683047	219545	311280	312640	
194671	2	-122.355317	47.678734	219546	309514	310794	
194672	1	-122.289360	47.611017	219547	308220	309500	

	REPORTNO	STATUS	ADDRTYPE	INTKEY	...	ROADCOND	\
0	3502005	Matched	Intersection	37475.0	...	Wet	
1	2607959	Matched	Block	NaN	...	Wet	
2	1482393	Matched	Block	NaN	...	Dry	
3	3503937	Matched	Block	NaN	...	Dry	
4	1807429	Matched	Intersection	34387.0	...	Wet	
...	
194668	E871089	Matched	Block	NaN	...	Dry	
194669	E876731	Matched	Block	NaN	...	Wet	
194670	3809984	Matched	Intersection	24760.0	...	Dry	
194671	3810083	Matched	Intersection	24349.0	...	Dry	
194672	E868008	Matched	Block	NaN	...	Wet	

	LIGHTCOND	PEDROWNOTGRNT	SDOTCOLNUM	SPEEDING	ST_COLCODE	\
0	Daylight	NaN	NaN	NaN	10	
1	Dark - Street Lights On	NaN	6354039.0	NaN	11	
2	Daylight	NaN	4323031.0	NaN	32	
3	Daylight	NaN	NaN	NaN	23	
4	Daylight	NaN	4028032.0	NaN	10	
...	
194668	Daylight	NaN	NaN	NaN	24	
194669	Daylight	NaN	NaN	NaN	13	
194670	Daylight	NaN	NaN	NaN	28	
194671	Dusk	NaN	NaN	NaN	5	
194672	Daylight	NaN	NaN	NaN	14	

	ST_COLDESC	SEGLANEKEY	\
0	Entering at angle	0	
1	From same direction - both going straight - bo...	0	

2	One parked--one moving	0
3	From same direction - all others	0
4	Entering at angle	0
...
194668	From opposite direction - both moving - head-on	0
194669	From same direction - both going straight - bo...	0
194670	From opposite direction - one left turn - one ...	0
194671	Vehicle Strikes Pedalcyclist	4308
194672	From same direction - both going straight - on...	0

	CROSSWALKKEY	HITPARKEDCAR
0	0	N
1	0	N
2	0	N
3	0	N
4	0	N
...
194668	0	N
194669	0	N
194670	0	N
194671	0	N
194672	0	N

[194673 rows x 38 columns]

Now we will drop all the columns that are not required & cleanup the data

```
[9]: data1 = data.drop(columns = ['OBJECTID', 'SEVERITYCODE.1', 'REPORTNO', 'INCKEY', 'COLDEKEY',
                                'X', 'Y', 'STATUS', 'ADDRTYPE',
                                'INTKEY', 'LOCATION', 'EXCEPTRSNCODE',
                                'EXCEPTRSNDESC', 'SEVERITYDESC', 'INCDATE',
                                'INCDTTM', 'JUNCTIONTYPE', 'SDOT_COLCODE',
                                'SDOT_COLDESC', 'PEDROWNOTGRNT', 'SDOTCOLNUM',
                                'ST_COLCODE', 'ST_COLDESC', 'SEGLANEKEY',
                                'CROSSWALKKEY', 'HITPARKEDCAR', 'PEDCOUNT', 'PEDCYLCOUNT',
                                'PERSONCOUNT', 'VEHCOUNT', 'COLLISIONTYPE',
                                'SPEEDING', 'UNDERINFL', 'INATTENTIONIND'])
```

```
[14]: data1
```

	SEVERITYCODE	WEATHER	ROADCOND	LIGHTCOND
0	2	Overcast	Wet	Daylight
1	1	Raining	Wet	Dark - Street Lights On
2	1	Overcast	Dry	Daylight
3	1	Clear	Dry	Daylight
4	2	Raining	Wet	Daylight

```

...
194668      2      Clear      Dry      Daylight
194669      1      Raining     Wet      Daylight
194670      2      Clear      Dry      Daylight
194671      2      Clear      Dry      Dusk
194672      1      Clear      Wet      Daylight

```

```
[194673 rows x 4 columns]
```

In its original form, this data is not fit for analysis. For one, there are many columns that we will not use for this model. Also, most of the features are of type object, when they should be numerical type.

We must use label encoding to convert the features to our desired data type:

```
[17]: data1["WEATHER"] = data1["WEATHER"].astype('category')
data1["ROADCOND"] = data1["ROADCOND"].astype('category')
data1["LIGHTCOND"] = data1["LIGHTCOND"].astype('category')

data1["WEATHER_CAT"] = data1["WEATHER"].cat.codes
data1["ROADCOND_CAT"] = data1["ROADCOND"].cat.codes
data1["LIGHTCOND_CAT"] = data1["LIGHTCOND"].cat.codes
```

```
[20]: data1.head()
```

```
[20]: SEVERITYCODE  WEATHER ROADCOND      LIGHTCOND  WEATHER_CAT  \
0              2  Overcast      Wet      Daylight           4
1              1   Raining     Wet  Dark - Street Lights On      6
2              1  Overcast      Dry      Daylight           4
3              1    Clear      Dry      Daylight           1
4              2   Raining     Wet      Daylight           6

      ROADCOND_CAT  LIGHTCOND_CAT
0              8              5
1              8              2
2              0              5
3              0              5
4              8              5
```

```
[22]: data1.dtypes
```

```
[22]: SEVERITYCODE      int64
WEATHER      category
ROADCOND      category
LIGHTCOND      category
WEATHER_CAT      int8
ROADCOND_CAT      int8
LIGHTCOND_CAT      int8
```

dtype: object

```
[24]: data1.columns
```

```
[24]: Index(['SEVERITYCODE', 'WEATHER', 'ROADCOND', 'LIGHTCOND', 'WEATHER_CAT',  
        'ROADCOND_CAT', 'LIGHTCOND_CAT'],  
        dtype='object')
```

2.1 Analyzing Value Counts

```
[27]: data1["SEVERITYCODE"].value_counts()
```

```
[27]: 1    136485  
      2     58188  
      Name: SEVERITYCODE, dtype: int64
```

```
[28]: data1["WEATHER"].value_counts()
```

```
[28]: Clear                111135  
      Raining             33145  
      Overcast            27714  
      Unknown             15091  
      Snowing              907  
      Other                832  
      Fog/Smog/Smoke       569  
      Sleet/Hail/Freezing Rain 113  
      Blowing Sand/Dirt     56  
      Severe Crosswind     25  
      Partly Cloudy         5  
      Name: WEATHER, dtype: int64
```

```
[29]: data1["ROADCOND"].value_counts()
```

```
[29]: Dry                124510  
      Wet               47474  
      Unknown           15078  
      Ice               1209  
      Snow/Slush        1004  
      Other             132  
      Standing Water    115  
      Sand/Mud/Dirt      75  
      Oil                64  
      Name: ROADCOND, dtype: int64
```

```
[30]: data1["LIGHTCOND"].value_counts()
```

```
[30]: Daylight          116137
      Dark - Street Lights On    48507
      Unknown              13473
      Dusk                  5902
      Dawn                  2502
      Dark - No Street Lights    1537
      Dark - Street Lights Off   1199
      Other                  235
      Dark - Unknown Lighting    11
      Name: LIGHTCOND, dtype: int64
```

Our target variable SEVERITYCODE is only 42% balanced. In fact, severitycode in class 1 is nearly three times the size of class 2.

We can fix this by downsampling the majority class:

```
[33]: data1_majority = data1[data1.SEVERITYCODE==1]
      data1_minority = data1[data1.SEVERITYCODE==2]

      #Downsample majority class
      data1_majority_downsampled = resample(data1_majority,
                                             replace=False,
                                             n_samples=58188,
                                             random_state=123)

      data1_balanced = pd.concat([data1_majority_downsampled, data1_minority])

      data1_balanced.SEVERITYCODE.value_counts()
```

```
[33]: 2    58188
      1    58188
      Name: SEVERITYCODE, dtype: int64
```

Now the data is ready to be analyzed.

3 Methodology

Our data is now ready to be fed into machine learning models.

We will use the following models:

K-Nearest Neighbor (KNN) KNN will help us predict the severity code of an outcome by finding the most similar to data point within k distance.

Decision Tree A decision tree model gives us a layout of all possible outcomes so we can fully analyze the consequences of a decision. In context, the decision tree observes all possible outcomes of different weather conditions.

Logistic Regression Because our dataset only provides us with two severity code outcomes, our model will only predict one of those two classes. This makes our data binary, which is perfect to use with logistic regression.

4 Initialization

Define x and y:

```
[35]: x = np.asarray(data1_balanced[['WEATHER_CAT', 'ROADCOND_CAT', 'LIGHTCOND_CAT']])  
      x[0:5]
```

```
[35]: array([[ 6,  8,  2],  
            [ 1,  0,  5],  
            [10,  7,  8],  
            [ 1,  0,  5],  
            [ 1,  0,  5]], dtype=int8)
```

```
[36]: y = np.asarray(data1_balanced['SEVERITYCODE'])  
      y [0:5]
```

```
[36]: array([1, 1, 1, 1, 1], dtype=int64)
```

Normalize the dataset

```
[39]: x = preprocessing.StandardScaler().fit(x).transform(x)  
      x[0:5]
```

```
[39]: array([[ 1.15236718,  1.52797946, -1.21648407],  
            [-0.67488    , -0.67084969,  0.42978835],  
            [ 2.61416492,  1.25312582,  2.07606076],  
            [-0.67488    , -0.67084969,  0.42978835],  
            [-0.67488    , -0.67084969,  0.42978835]])
```

Train/Test Split We will use 30% of our data for testing and 70% for training:

```
[41]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3,  
    ↪random_state=4)  
      print ('Train set:', x_train.shape, y_train.shape)  
      print ('Test set:', x_test.shape, y_test.shape)
```

Train set: (81463, 3) (81463,)

Test set: (34913, 3) (34913,)

K-Nearest Neighbors (KNN)

```
[43]: k = 25
```

```
[45]: neigh = KNeighborsClassifier(n_neighbors = k).fit(x_train,y_train)
      neigh

      Kyhat = neigh.predict(x_test)
      Kyhat[0:5]
```

```
[45]: array([2, 2, 1, 1, 2], dtype=int64)
```

Decision Tree

```
[47]: # Building the Decision Tree
      data1_Tree = DecisionTreeClassifier(criterion="entropy", max_depth = 7)
      data1_Tree
      data1_Tree.fit(x_train,y_train)
```

```
[47]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=7,
                             max_features=None, max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, presort=False,
                             random_state=None, splitter='best')
```

```
[51]: # Train Model & Predict
      predTree = data1_Tree.predict(x_test)
      print (predTree [0:5])
      print (y_test [0:5])
```

```
[2 2 1 1 2]
```

```
[2 2 1 1 1]
```

Logistic Regression

```
[54]: # Building the LR Model
      LR = LogisticRegression(C=6, solver='liblinear').fit(x_train,y_train)
      LR
```

```
[54]: LogisticRegression(C=6, class_weight=None, dual=False, fit_intercept=True,
                          intercept_scaling=1, l1_ratio=None, max_iter=100,
                          multi_class='warn', n_jobs=None, penalty='l2',
                          random_state=None, solver='liblinear', tol=0.0001, verbose=0,
                          warm_start=False)
```

```
[55]: # Train Model & Predict
      LRYhat = LR.predict(x_test)
      LRYhat
```

```
[55]: array([1, 2, 1, ..., 2, 2, 2], dtype=int64)
```



```
[56]: yhat_prob = LR.predict_proba(x_test)
      yhat_prob
```

```
[56]: array([[0.57295252, 0.42704748],
          [0.47065071, 0.52934929],
          [0.67630201, 0.32369799],
          ...,
          [0.46929132, 0.53070868],
          [0.47065071, 0.52934929],
          [0.46929132, 0.53070868]])
```

5 Results & Evaluation

Now we will check the accuracy of our models:

K-Nearest Neighbor

```
[59]: # Jaccard Similarity Score
      jaccard_similarity_score(y_test, Kyhat)
```

```
C:\Users\avira\Anaconda3\lib\site-
packages\sklearn\metrics\classification.py:635: DeprecationWarning:
jaccard_similarity_score has been deprecated and replaced with jaccard_score. It
will be removed in version 0.23. This implementation has surprising behavior for
binary and multiclass classification tasks.
'and multiclass classification tasks.', DeprecationWarning)
```

```
[59]: 0.564001947698565
```

```
[60]: # F1-SCORE
      f1_score(y_test, Kyhat, average='macro')
```

```
[60]: 0.5401775308974308
```

Model is most accurate when k is 25.

Decision Tree

```
[61]: # Jaccard Similarity Score
      jaccard_similarity_score(y_test, DTyhat)
```

```
C:\Users\avira\Anaconda3\lib\site-
packages\sklearn\metrics\classification.py:635: DeprecationWarning:
jaccard_similarity_score has been deprecated and replaced with jaccard_score. It
will be removed in version 0.23. This implementation has surprising behavior for
binary and multiclass classification tasks.
'and multiclass classification tasks.', DeprecationWarning)
```

```
[61]: 0.5664365709048206
```

```
[64]: # F1-SCORE  
f1_score(y_test, DTyhat, average='macro')
```

```
[64]: 0.5450597937389444
```

Model is most accurate with a max depth of 7.

Logistic Regression

```
[66]: # Jaccard Similarity Score  
jaccard_similarity_score(y_test, LRyhat)
```

```
C:\Users\avira\Anaconda3\lib\site-  
packages\sklearn\metrics\classification.py:635: DeprecationWarning:  
jaccard_similarity_score has been deprecated and replaced with jaccard_score. It  
will be removed in version 0.23. This implementation has surprising behavior for  
binary and multiclass classification tasks.  
'and multiclass classification tasks.', DeprecationWarning)
```

```
[66]: 0.5260218256809784
```

```
[67]: # F1-SCORE  
f1_score(y_test, LRyhat, average='macro')
```

```
[67]: 0.511602093963383
```

```
[68]: ##### LOGLOSS  
yhat_prob = LR.predict_proba(x_test)  
log_loss(y_test, yhat_prob)
```

```
[68]: 0.6849535383198887
```

Model is most accurate when hyperparameter C is 6.

6 Discussion

In the beginning of this notebook, we had categorical data that was of type 'object'. This is not a data type that we could have fed through an algorithm, so label encoding was used to create new classes that were of type int8; a numerical data type.

After solving that issue we were presented with another - imbalanced data. As mentioned earlier, class 1 was nearly three times larger than class 2. The solution to this was downsampling the majority class with sklearn's resample tool. We downsampled to match the minority class exactly with 58188 values each.

Once we analyzed and cleaned the data, it was then fed through three ML models; K-Nearest Neighbor, Decision Tree and Logistic Regression. Although the first two are ideal for this project, logistic regression made the most sense because of its binary nature.

Evaluation metrics used to test the accuracy of our models were jaccard index, f-1 score and logloss for logistic regression. Choosing different k, max depth and hyperparameter C values helped to improve our accuracy to be the best possible.

7 Conclusion

Based on historical data from weather conditions pointing to certain classes, we can conclude that particular weather conditions have a somewhat impact on whether or not travel could result in property damage (class 1) or injury (class 2).